

notebook_Mask_RCNN_car_damage_pred_nasir

August 22, 2019

0.1 DamageDetection with MaskRCNN

- Nasir Uddin
- Data visualization of car damage images and automated car damage detection example.

```
[1]: #import all the packages including custom functions of Matterport Mask R-CNN  
    → repository  
import os  
import sys  
import itertools  
import math  
import logging  
import json  
import re  
import random  
from collections import OrderedDict  
import numpy as np  
import matplotlib  
import matplotlib.pyplot as plt  
import matplotlib.patches as patches  
import matplotlib.lines as lines  
from matplotlib.patches import Polygon  
  
# Import Mask RCNN  
#sys.path.append(ROOT_DIR) # To find local version of the library  
from mrcnn import utils  
from mrcnn import visualize  
from mrcnn.visualize import display_images  
from mrcnn import model  
import mrcnn.model as modellib  
from mrcnn.model import log  
import cv2  
import custom, custom_1  
import imgaug, h5py, IPython  
  
%matplotlib inline
```

Using TensorFlow backend.

Setting up the configuration - root directory, data path setting up the log file path and model object(weight matrix) for inference (prediction)

```
[2]: # Root directory of the project
ROOT_DIR = os.getcwd()
sys.path.append(ROOT_DIR) # To find local version of the library
MODEL_DIR = os.path.join(ROOT_DIR, "logs")
custom_WEIGHTS_PATH = "mask_rcnn_coco.h5" # TODO: update this path for best_
      ↪performing iteration weights
config = custom.CustomConfig()
custom_DIR = os.path.join(ROOT_DIR, "custom/")
custom_DIR
```

```
[2]: '/home/nasir/Desktop/carcnn/car-damage-detection-using-CNN/custom/'
```

loading the data

```
[3]: # Load dataset
dataset = custom_1.CustomDataset()
dataset.load_custom(custom_DIR, "train")

# Must call before using the dataset
dataset.prepare()

print("Image Count: {}".format(len(dataset.image_ids)))
print("Class Count: {}".format(dataset.num_classes))
for i, info in enumerate(dataset.class_info):
    print("{:3}. {:50}".format(i, info['name']))
```

Image Count: 49

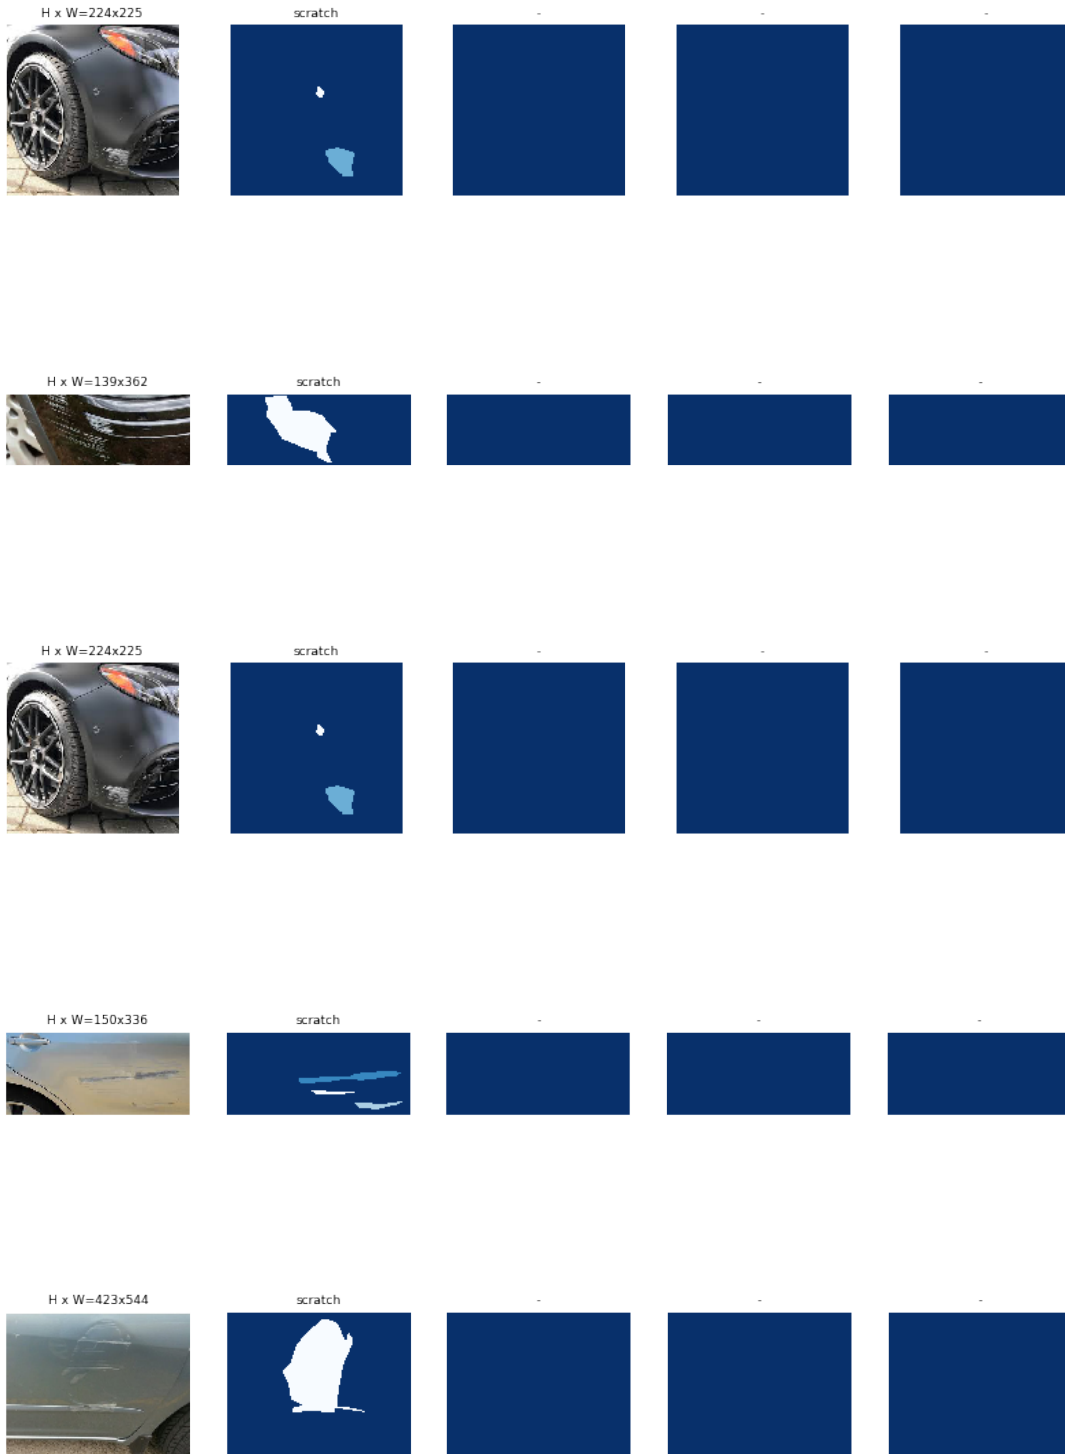
Class Count: 2

0. BG

1. scratch

We will visualize few car damage(scratch) images

```
[4]: # Load and display random samples
image_ids = np.random.choice(dataset.image_ids, 5)
for image_id in image_ids:
    image = dataset.load_image(image_id)
    mask, class_ids = dataset.load_mask(image_id)
    visualize.display_top_masks(image, mask, class_ids, dataset.class_names)
```



Make Bounding Box(BB)with annotated damage mask for a typical car image.

```
[5]: image_id = random.choice(dataset.image_ids)
     image = dataset.load_image(image_id)
```

```

mask, class_ids = dataset.load_mask(image_id)
# Compute Bounding box
bbox = utils.extract_bboxes(mask)

# Display image and additional stats
print("image_id ", image_id, dataset.image_reference(image_id))
log("image", image)
log("mask", mask)
log("class_ids", class_ids)
log("bbox", bbox)
# Display image and instances
visualize.display_instances(image, bbox, mask, class_ids, dataset.class_names)

```

```

image_id 41 /home/nasir/Desktop/carcnn/car-damage-detection-using-
CNN/custom/train/image46.png
image          shape: (528, 705, 3)          min:    0.00000  max:
255.00000  uint8
mask          shape: (528, 705, 1)          min:    0.00000  max:
1.00000  bool
class_ids     shape: (1,)                   min:    1.00000  max:
1.00000  int32
bbox          shape: (1, 4)                 min:   184.00000  max:
478.00000  int32

```



We see some the components of image annotations. Mainly it has x and y co-ordinate of all labeled damages('polygon') and class name(here 'scratch') for respective car image.

```
[7]: #Annotation file load
annotations1 = json.load(open(os.path.join(ROOT_DIR, "custom/train/
→via_region_data.json"), encoding="utf8"))
annotations = list(annotations1.values())
annotations = [a for a in annotations if a['regions']]
annotations[0]
```

```
[7]: {'fileref': '',
      'size': 46041,
      'filename': 'image2.jpg',
      'base64_img_data': '',
      'file_attributes': {},
      'regions': {'0': {'shape_attributes': {'name': 'polygon',
      'all_points_x': [428,
      429,
      480,
      518,
      557,
      577,
      610,
      660,
      642,
      578,
      579,
      585,
      590,
      574,
      580,
      516,
      507,
      474,
      427,
      426,
      412,
      412,
      430,
      470,
      452,
      428],
      'all_points_y': [232,
      216,
      198,
      193,
```

```

212,
238,
237,
242,
248,
248,
260,
292,
343,
409,
417,
441,
443,
427,
413,
381,
324,
301,
288,
249,
231,
232]},
'region_attributes': {'Scratch': 'scratch'}},
'1': {'shape_attributes': {'name': 'polygon',
'all_points_x': [470, 500, 578, 718, 670, 594, 553, 510, 469, 448, 470],
'all_points_y': [516, 548, 562, 557, 569, 595, 587, 600, 576, 552, 516]},
'region_attributes': {'Scratch': 'scratch'}}}]

```

If we have to quantify a car damage, we need to know the x and y coordinates of the polygon to calculate area of the marked/detected damage. This is for 2nd damage polygon of 'image2.jpg'

```

[8]: annotations[1]['regions']['0']['shape_attributes']
l = []
for d in annotations[1]['regions']['0']['shape_attributes'].values():
    l.append(d)
display('x co-ordinates of the damage:', l[1])
display('y co-ordinates of the damage:', l[2])

```

'x co-ordinates of the damage:'

[293, 360, 349, 308, 293]

'y co-ordinates of the damage:'

[303, 330, 314, 302, 303]

For prediction or damage detection we need to use the model as inference mode. Model description consists of important model information like CNN architecture name('resnet101'), ROI threshold(0.9 as defined), configuration description, weightage of different loss components, mask shape, WEIGHT_DECAY etc.

```
[9]: config = custom.CustomConfig()
ROOT_DIR = "/home/nasir/Desktop/carcnn/car-damage-detection-using-CNN"
CUSTOM_DIR = os.path.join(ROOT_DIR + "/custom/")
print(CUSTOM_DIR)
class InferenceConfig(config.__class__):
    # Run detection on one image at a time
    GPU_COUNT = 1
    IMAGES_PER_GPU = 1

config = InferenceConfig()
config.display()

# Device to load the neural network on.
# Useful if you're training a model on the same
# machine, in which case use CPU and leave the
# GPU for training.
DEVICE = "/cpu:0" # /cpu:0 or /gpu:0

# Inspect the model in training or inference modes
# values: 'inference' or 'training'
# TODO: code for 'training' test mode not ready yet
TEST_MODE = "inference"
```

/home/nasir/Desktop/carcnn/car-damage-detection-using-CNN/custom/

Configurations:

BACKBONE	resnet101
BACKBONE_STRIDES	[4, 8, 16, 32, 64]
BATCH_SIZE	1
BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
COMPUTE_BACKBONE_SHAPE	None
DETECTION_MAX_INSTANCES	100
DETECTION_MIN_CONFIDENCE	0.9
DETECTION_NMS_THRESHOLD	0.3
FPN_CLASSIF_FC_LAYERS_SIZE	1024
GPU_COUNT	1
GRADIENT_CLIP_NORM	5.0
IMAGES_PER_GPU	1
IMAGE_CHANNEL_COUNT	3
IMAGE_MAX_DIM	1024
IMAGE_META_SIZE	14
IMAGE_MIN_DIM	800
IMAGE_MIN_SCALE	0

```

IMAGE_RESIZE_MODE      square
IMAGE_SHAPE            [1024 1024    3]
LEARNING_MOMENTUM      0.9
LEARNING_RATE          0.001
LOSS_WEIGHTS           {'rpn_class_loss': 1.0, 'rpn_bbox_loss': 1.0,
                        'mrcnn_class_loss': 1.0, 'mrcnn_bbox_loss': 1.0, 'mrcnn_mask_loss': 1.0}
MASK_POOL_SIZE         14
MASK_SHAPE             [28, 28]
MAX_GT_INSTANCES       100
MEAN_PIXEL             [123.7 116.8 103.9]
MINI_MASK_SHAPE        (56, 56)
NAME                   damage
NUM_CLASSES            2
POOL_SIZE              7
POST_NMS_ROIS_INFERENCE 1000
POST_NMS_ROIS_TRAINING  2000
PRE_NMS_LIMIT          6000
ROI_POSITIVE_RATIO     0.33
RPN_ANCHOR_RATIOS      [0.5, 1, 2]
RPN_ANCHOR_SCALES      (32, 64, 128, 256, 512)
RPN_ANCHOR_STRIDE      1
RPN_BBOX_STD_DEV       [0.1 0.1 0.2 0.2]
RPN_NMS_THRESHOLD      0.7
RPN_TRAIN_ANCHORS_PER_IMAGE 256
STEPS_PER_EPOCH        100
TOP_DOWN_PYRAMID_SIZE  256
TRAIN_BN               False
TRAIN_ROIS_PER_IMAGE   200
USE_MINI_MASK          True
USE_RPN_ROIS           True
VALIDATION_STEPS       50
WEIGHT_DECAY           0.0001

```

Helper function to visualize predicted damage masks and loading the model weights for prediction

```

[11]: def get_ax(rows=1, cols=1, size=16):
        """Return a Matplotlib Axes array to be used in
        all visualizations in the notebook. Provide a
        central point to control graph sizes.

        Adjust the size attribute to control how big to render images
        """
        _, ax = plt.subplots(rows, cols, figsize=(size*cols, size*rows))
        return ax

```



```

from importlib import reload # was constantly changin the visualization, so I
    ↳decided to reload it instead of notebook
reload(visualize)

# Create model in inference mode
import tensorflow as tf
with tf.device(DEVICE):
    model = modellib.MaskRCNN(mode="inference", model_dir=MODEL_DIR,
                              config=config)

# load the last best model you trained
# weights_path = model.find_last()[1]
custom_WEIGHTS_PATH = '/home/nasir/Desktop/carcnn/
    ↳car-damage-detection-using-CNN/logs/scratch20190822T1650/
    ↳mask_rcnn_scratch_0015.h5'
# Load weights
print("Loading weights ", custom_WEIGHTS_PATH)
model.load_weights(custom_WEIGHTS_PATH, by_name=True)

```

Loading weights /home/nasir/Desktop/carcnn/car-damage-detection-using-CNN/logs/scratch20190822T1650/mask_rcnn_scratch_0015.h5
 Re-starting from epoch 15

Loading validation data-set for prediction

```

[12]: dataset = custom_1.CustomDataset()
dataset.load_custom(CUSTOM_DIR, 'val')
dataset.prepare()
print('Images: {} \n classes: {}'.format(len(dataset.image_ids), dataset.
    ↳class_names))

```

Images: 6
 classes: ['BG', 'scratch']

Visualize model weight matrix descriptive statistics(shapes, histograms)

```

[13]: visualize.display_weight_stats(model)

```

<IPython.core.display.HTML object>

Prediction on a random validation image(image53.jpeg)

```

[14]: image_id = random.choice(dataset.image_ids)
image, image_meta, gt_class_id, gt_bbox, gt_mask =\
    modellib.load_image_gt(dataset, config, image_id, use_mini_mask=False)
info = dataset.image_info[image_id]
print("image ID: {}.{} ({}). {}".format(info["source"], info["id"], image_id,

```

```

dataset.image_reference(image_id))

# Run object detection
results = model.detect([image], verbose=1)

# Display results
ax = get_ax(1)
r = results[0]
visualize.display_instances(image, r['rois'], r['masks'], r['class_ids'],
                           dataset.class_names, r['scores'], ax=ax,
                           title="Predictions")
log("gt_class_id", gt_class_id)
log("gt_bbox", gt_bbox)
log("gt_mask", gt_mask)
print('The car has: {} damages'.format(len(dataset.
→image_info[image_id]['polygons'])))


```

image ID: scratch.image54.jpg (3) /home/nasir/Desktop/carcnn/car-damage-detection-using-CNN/custom/val/image54.jpg

Processing 1 images

image	shape: (1024, 1024, 3)	min: 0.00000	max:
255.00000 uint8			
molded_images	shape: (1, 1024, 1024, 3)	min: -123.70000	max:
151.10000 float64			
image metas	shape: (1, 14)	min: 0.00000	max:
1024.00000 int64			
anchors	shape: (1, 261888, 4)	min: -0.35390	max:
1.29134 float32			
gt_class_id	shape: (1,)	min: 1.00000	max:
1.00000 int32			
gt_bbox	shape: (1, 4)	min: 221.00000	max:
856.00000 int32			
gt_mask	shape: (1024, 1024, 1)	min: 0.00000	max:
1.00000 bool			

The car has:1 damages



On another image

```
[15]: image_id = random.choice(dataset.image_ids)
image, image_meta, gt_class_id, gt_bbox, gt_mask = \
    modellib.load_image_gt(dataset, config, image_id, use_mini_mask=False)
info = dataset.image_info[image_id]
print("image ID: {}.{} ({}). {}".format(info["source"], info["id"], image_id,
                                         dataset.image_reference(image_id)))

# Run object detection
results = model.detect([image], verbose=1)

# Display results
```

```

ax = get_ax(1)
r = results[0]
visualize.display_instances(image, r['rois'], r['masks'], r['class_ids'],
                           dataset.class_names, r['scores'], ax=ax,
                           title="Predictions")
log("gt_class_id", gt_class_id)
log("gt_bbox", gt_bbox)
log("gt_mask", gt_mask)
print('The car has: {} damages'.format(len(dataset.
→image_info[image_id]['polygons'])))


```

image ID: scratch.image51.png (0) /home/nasir/Desktop/carcnn/car-damage-detection-using-CNN/custom/val/image51.png

Processing 1 images

image	shape: (1024, 1024, 3)	min: 0.00000	max:
255.00000 uint8			
molded_images	shape: (1, 1024, 1024, 3)	min: -123.70000	max:
151.10000 float64			
image metas	shape: (1, 14)	min: 0.00000	max:
1024.00000 int64			
anchors	shape: (1, 261888, 4)	min: -0.35390	max:
1.29134 float32			
gt_class_id	shape: (1,)	min: 1.00000	max:
1.00000 int32			
gt_bbox	shape: (1, 4)	min: 351.00000	max:
689.00000 int32			
gt_mask	shape: (1024, 1024, 1)	min: 0.00000	max:
1.00000 bool			
The car has:1 damages			

Predictions



Good prediction considering training with only 49 images and 15 epochs.
Thanks!

[]: